


1. Download the data from [here](#). You have to use data.csv file for this assignment
2. Code the model to classify data like below image. You can use any number of units in your Dense layers.



```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
import pandas as pd
import numpy as np
```

```
path = "/content/drive/MyDrive/Assignment_data/data.csv"
df = pd.read_csv(path)

df.head()
```

	f1	f2	label	
0	0.450564	1.074305	0.0	
1	0.085632	0.967682	0.0	
2	0.117326	0.971521	1.0	
3	0.982179	-0.380408	0.0	
4	-0.720352	0.955850	0.0	

```
from sklearn.model_selection import train_test_split
y = df['label'].values
X = df.drop(['label'], axis=1)
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.20, stratify=y)
X_train.shape, X_test.shape, Y_train.shape, Y_test.shape

((16000, 2), (4000, 2), (16000,), (4000,))
```

▼ 3. Writing Callbacks

You have to implement the following callbacks

- Write your own callback function, that has to print the micro F1 score and AUC score after each epoch. Do not use `tf.keras.metrics` for calculating AUC and F1 score.
- Save your model at every epoch if your validation accuracy is improved from previous epoch.
- You have to decay learning based on below conditions

Cond1. If your validation accuracy at that epoch is less than previous epoch and learning rate by 10%.

Cond2. For every 3rd epoch, decay your learning rate by 5%.

- If you are getting any NaN values (either weights or loss) while training, you have to terminate your training.
- You have to stop the training if your validation accuracy is not increased in last 2 epochs.
- Use tensorboard for every model and analyse your scalar plots and histograms. (you need to upload the screenshots and write the observations for each model for evaluation)

Model-1

1. Use tanh as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

```
import tensorflow as tf
from tensorflow.keras.layers import Dense, Input, Activation
from tensorflow.keras.models import Model
import random as rn

from sklearn.metrics import f1_score
from sklearn.metrics import roc_auc_score

class metric_callback(tf.keras.callbacks.Callback):
    def __init__(self, validation_data):
        self.x_test = validation_data[0]
        self.y_test = validation_data[1]
    def on_train_begin(self, logs={}):
        ## on begin of training, we are creating a instance variable called history
        ## it is a dict with keys [loss, acc, val_loss, val_acc]
        self.history = {'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': [], 'val_f1s': []}
    def on_epoch_end(self, epoch, logs={}):
        true_positives = 0
        ## on end of each epoch, we will get logs and update the self.history dict
        self.history['loss'].append(logs.get('loss'))
        self.history['accuracy'].append(logs.get('accuracy'))

        if logs.get('val_loss', -1) != -1:
            self.history['val_loss'].append(logs.get('val_loss'))
        if logs.get('val_accuracy', -1) != -1:
            self.history['val_accuracy'].append(logs.get('val_accuracy'))
        # we can get a list of all predicted values at the end of the epoch
        # we can use these predicted value and the true values to calculate any custom eva
        # Here we are taking log of all true positives and then taking average of it
        y_pred = self.model.predict(self.x_test)
        y_label_pred = np.argmax(y_pred, axis=1)

        # we can also calculate predefined metrics such as precision, recall, etc. using cal
        f1score = round(f1_score(self.y_test, y_label_pred, average='micro'), 4)
        self.history['val_f1score'].append(f1score)
        auc = round(roc_auc_score(self.y_test, y_label_pred), 4)
        self.history['val_AUC'].append(auc)
```

```

class modelSave_callback(tf.keras.callbacks.Callback):
    def __init__(self, path):
        self.filepath_name = path

    def on_train_begin(self, logs={}):
        ## on begin of training, we are creating a instance variable called history
        ## it is a dict with keys [loss, acc, val_loss, val_acc]
        self.history={'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': [], 'val_f1s': []}

    def on_epoch_end(self, epoch, logs={}):
        true_positives=0
        ## on end of each epoch, we will get logs and update the self.history dict
        self.history['loss'].append(logs.get('loss'))
        self.history['accuracy'].append(logs.get('accuracy'))

        if logs.get('val_loss', -1) != -1:
            self.history['val_loss'].append(logs.get('val_loss'))
        if logs.get('val_accuracy', -1) != -1:
            self.history['val_accuracy'].append(logs.get('val_accuracy'))

        if epoch >= 2:
            if self.history.get('val_accuracy')[-1] > self.history.get('val_accuracy')[-2]:
                self.model.save(self.filepath_name)

class learning_rate_callback(tf.keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        ## on begin of training, we are creating a instance variable called history
        ## it is a dict with keys [loss, acc, val_loss, val_acc]
        self.history={'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': [], 'val_f1s': []}

    def on_epoch_end(self, epoch, logs={}):
        true_positives=0
        ## on end of each epoch, we will get logs and update the self.history dict
        self.history['loss'].append(logs.get('loss'))
        self.history['accuracy'].append(logs.get('accuracy'))

        if logs.get('val_loss', -1) != -1:
            self.history['val_loss'].append(logs.get('val_loss'))
        if logs.get('val_accuracy', -1) != -1:
            self.history['val_accuracy'].append(logs.get('val_accuracy'))

    def on_epoch_begin(self, epoch, logs={}):
        if epoch > 2:
            if self.history.get('val_accuracy')[-1] < self.history.get('val_accuracy')[-2] or
               if self.history.get('val_accuracy')[-1] < self.history.get('val_accuracy')[-2]:
                self.model.optimizer.lr = self.model.optimizer.lr*0.9
            else:
                self.model.optimizer.lr = 0.95*self.model.optimizer.lr

class terminateNaN_callback(tf.keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        ## on begin of training, we are creating a instance variable called history
        ## it is a dict with keys [loss, acc, val_loss, val_acc]
        self.history={'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': [], 'val_f1s': []}

```

```

def on_epoch_end(self, epoch, logs={}):
    true_positives=0
    ## on end of each epoch, we will get logs and update the self.history dict
    self.history['loss'].append(logs.get('loss'))
    self.history['accuracy'].append(logs.get('accuracy'))

    if logs.get('val_loss', -1) != -1:
        self.history['val_loss'].append(logs.get('val_loss'))
    if logs.get('val_accuracy', -1) != -1:
        self.history['val_accuracy'].append(logs.get('val_accuracy'))

    model_weights = self.model.get_weights()
    if model_weights is not None:
        if np.any([np.any(np.isnan(x)) for x in model_weights]):
            print("Invalid model weights and terminated at epoch {}".format(epoch))
            self.model.stop_training = True
    loss = logs.get('loss')
    if loss is not None:
        if np.isnan(loss) or np.isinf(loss):
            print("Invalid loss and terminated at epoch {}".format(epoch))
            self.model.stop_training = True

class modelcheck_callback(tf.keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        ## on begin of training, we are creating a instance variable called history
        ## it is a dict with keys [loss, acc, val_loss, val_acc]
        self.history={'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': [], 'val_f1s': []}

    def on_epoch_end(self, epoch, logs={}):
        true_positives=0
        ## on end of each epoch, we will get logs and update the self.history dict
        self.history['loss'].append(logs.get('loss'))
        self.history['accuracy'].append(logs.get('accuracy'))

        if logs.get('val_loss', -1) != -1:
            self.history['val_loss'].append(logs.get('val_loss'))
        if logs.get('val_accuracy', -1) != -1:
            self.history['val_accuracy'].append(logs.get('val_accuracy'))

        if epoch >= 3:
            if self.history.get('val_accuracy')[-1] <= self.history.get('val_accuracy')[-2]:
                print("Val_accuracy is not improved from last 2 epochs")
                self.model.stop_training = True

# there are other ways of doing this: https://www.dlology.com/blog/quick-guide-to-run-tens
%load_ext tensorboard
# Clear any logs from previous runs
!rm -rf ./logs/

from tensorflow.keras.callbacks import ModelCheckpoint, TerminateOnNaN, EarlyStopping, LearningRateScheduler

#Input layer

```

```

input_layer = Input(shape=(2,))
#Dense hidden layer
layer1 = Dense(32,activation='tanh',kernel_initializer=tf.keras.initializers.RandomUniform
layer2 = Dense(16,activation='tanh',kernel_initializer=tf.keras.initializers.RandomUniform
layer3 = Dense(8,activation='tanh',kernel_initializer=tf.keras.initializers.RandomUniform(
layer4 = Dense(4,activation='tanh',kernel_initializer=tf.keras.initializers.RandomUniform(
layer5 = Dense(2,activation='tanh',kernel_initializer=tf.keras.initializers.RandomUniform(
#output layer
output = Dense(1,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_norm
#Creating a model
model1 = Model(inputs=input_layer,outputs=output)

optimizer = tf.keras.optimizers.SGD(learning_rate=0.1)

model1.compile(optimizer=optimizer, loss='categorical_crossentropy',metrics=['accuracy'])


#Callbacks
mc=metric_callback(validation_data=[X_test,Y_test])
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
ms = modelSave_callback(filepath)
lr = learning_rate_callback()
tn = terminateNaN_callback()
mcheck = modelcheck_callback()
t_c = tf.keras.callbacks.TensorBoard(log_dir="./logs")

callback_lst = [mc,ms,lr,tn,mcheck,t_c]

model1.fit(X_train,Y_train,epochs=10,validation_data=(X_test,Y_test),batch_size=16,callbacks=callback_lst)

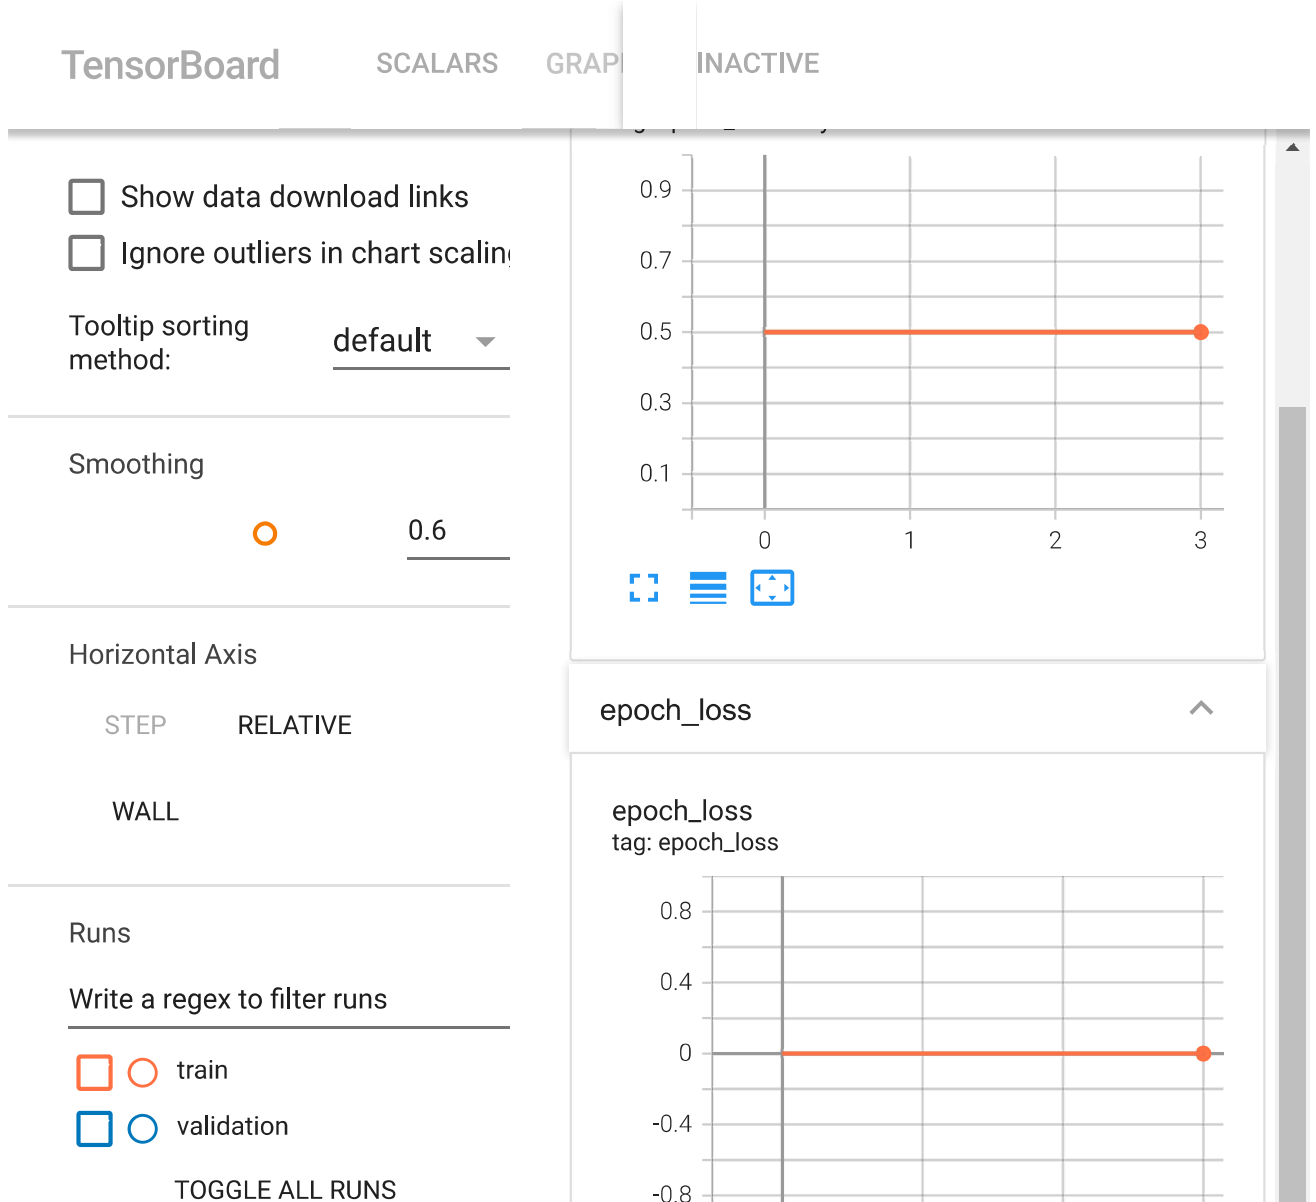
Epoch 1/10
 1/1000 [.....] - ETA: 6:54 - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 2/10
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 3/10
1000/1000 [=====] - 1s 1ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 4/10
 993/1000 [=====>.] - ETA: 0s - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
<keras.callbacks.History at 0x7fdf72564910>

```



```
%tensorboard --logdir logs
```

Reusing TensorBoard on port 6008 (pid 2411), started 0:06:30 ago. (Use '!kill 2411' to



1. Accuracy & AUC are same i.e. 0.5
2. Training terminated at epoch 4 because of no improvement in the "Val_Accuracy"
3. Both Train and test accuracies are same

evaluation accuracy vs iterations

Model-2

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

```

#Input layer
input_layer = Input(shape=(2,))
#Dense hidden layer
layer1 = Dense(32,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(
layer2 = Dense(16,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(
layer3 = Dense(8,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(
layer4 = Dense(4,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(
layer5 = Dense(2,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(
#output layer
output = Dense(1,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_norm
#Creating a model
model2 = Model(inputs=input_layer,outputs=output)

optimizer = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.9)

model2.compile(optimizer=optimizer, loss='categorical_crossentropy',metrics=['accuracy'])

#Callbacks
mc=metric_callback(validation_data=[X_test,Y_test])
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
ms = modelSave_callback(filepath)
lr = learning_rate_callback()
tn = terminateNaN_callback()
mcheck = modelcheck_callback()
t_c = tf.keras.callbacks.TensorBoard(log_dir="./logs")

callback_lst = [mc,ms,lr,tn,mcheck,t_c]

model2.fit(X_train,Y_train,epochs=10,validation_data=(X_test,Y_test),batch_size=16,callbacks=callback_lst)

Epoch 1/10
 1/1000 [.....] - ETA: 10:21 - loss: 0.0000e+00 - accuracy: 0.0000
981/1000 [=====>.] - ETA: 0s - loss: nan - accuracy: 0.4992
Invalid loss and terminated at epoch 0
1000/1000 [=====] - 2s 2ms/step - loss: nan - accuracy: 0.5000
<keras.callbacks.History at 0x7fdf72242b90>

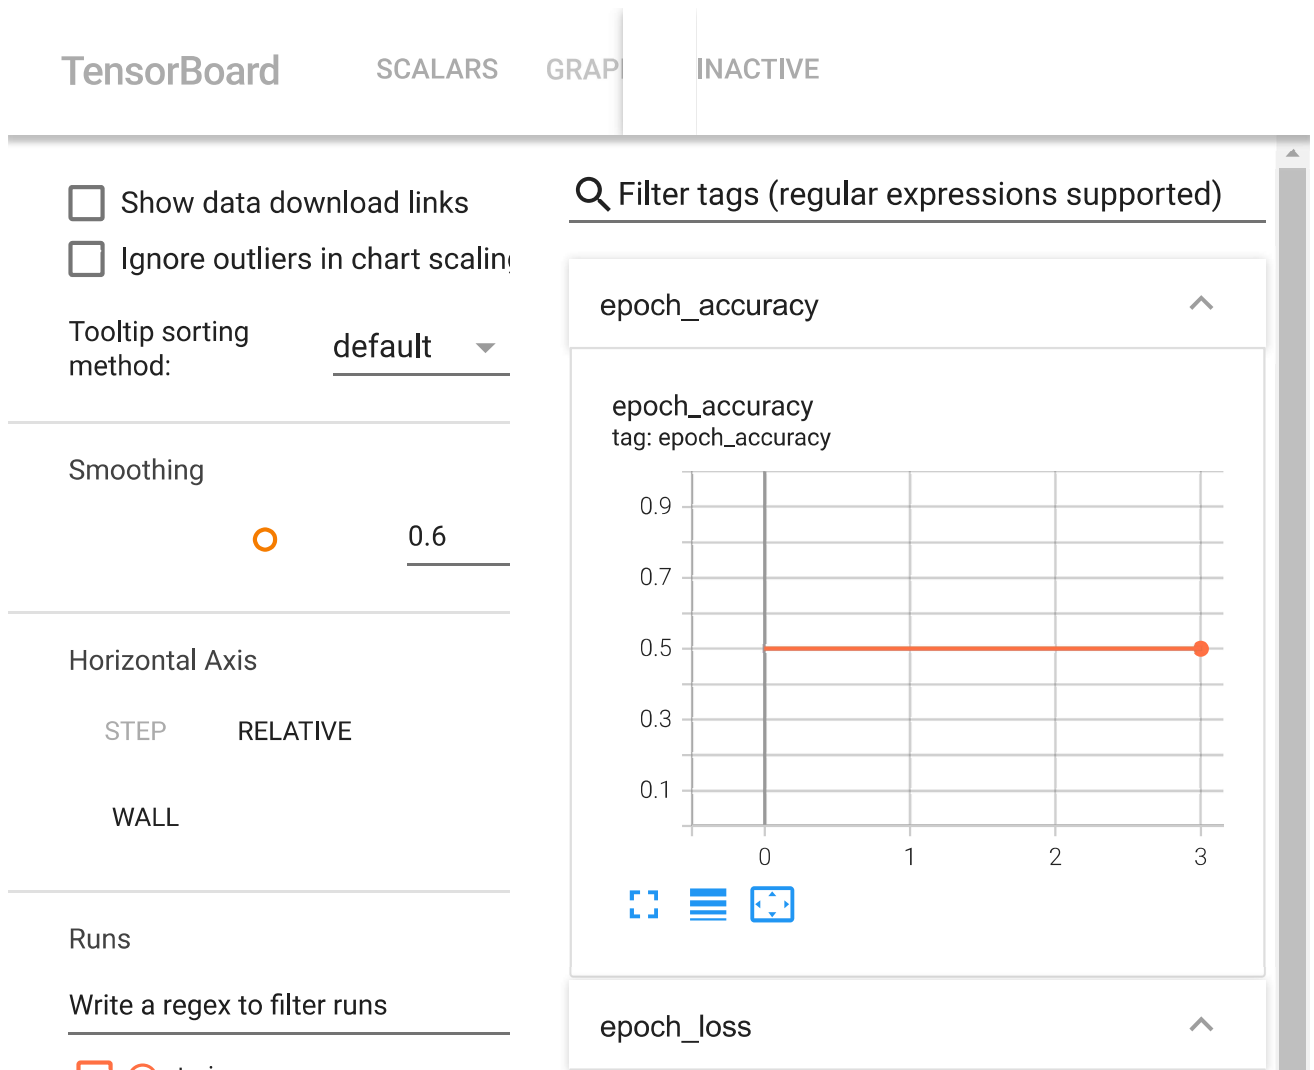
```



```

%tensorboard --logdir logs_model2

```

1. Accuracy & AUC are same i.e. 0.5
2. Training terminated at epoch 4 because of no improvement in the "Val_Accuracy"
3. Both Train and test accuracies are same

0.4

Model-3

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he_uniform() as initializer.
3. Analyze your output and training process.

```
#Input layer
input_layer = Input(shape=(2,))
#Dense hidden layer
layer1 = Dense(32,activation='relu',kernel_initializer=tf.keras.initializers.HeUniform(seed=1))
layer2 = Dense(16,activation='relu',kernel_initializer=tf.keras.initializers.HeUniform(seed=2))
layer3 = Dense(8,activation='relu',kernel_initializer=tf.keras.initializers.HeUniform(seed=3))
```

```

layer4 = Dense(4,activation='relu',kernel_initializer=tf.keras.initializers.HeUniform(seed
layer5 = Dense(2,activation='relu',kernel_initializer=tf.keras.initializers.HeUniform(seed
#output layer
output = Dense(1,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_norm
#Creating a model
model3 = Model(inputs=input_layer,outputs=output)

optimizer = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.9)

model3.compile(optimizer=optimizer, loss='categorical_crossentropy',metrics=['accuracy'])

#Callbacks
mc=metric_callback(validation_data=[X_test,Y_test])
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
ms = modelSave_callback(filepath)
lr = learning_rate_callback()
tn = terminateNaN_callback()
mcheck = modelcheck_callback()
t_c = tf.keras.callbacks.TensorBoard(log_dir="./logs")

callback_lst = [mc,ms,lr,tn,mcheck,t_c]


model3.fit(X_train,Y_train,epochs=10,validation_data=(X_test,Y_test),batch_size=16,callbacks=callback_lst)

```

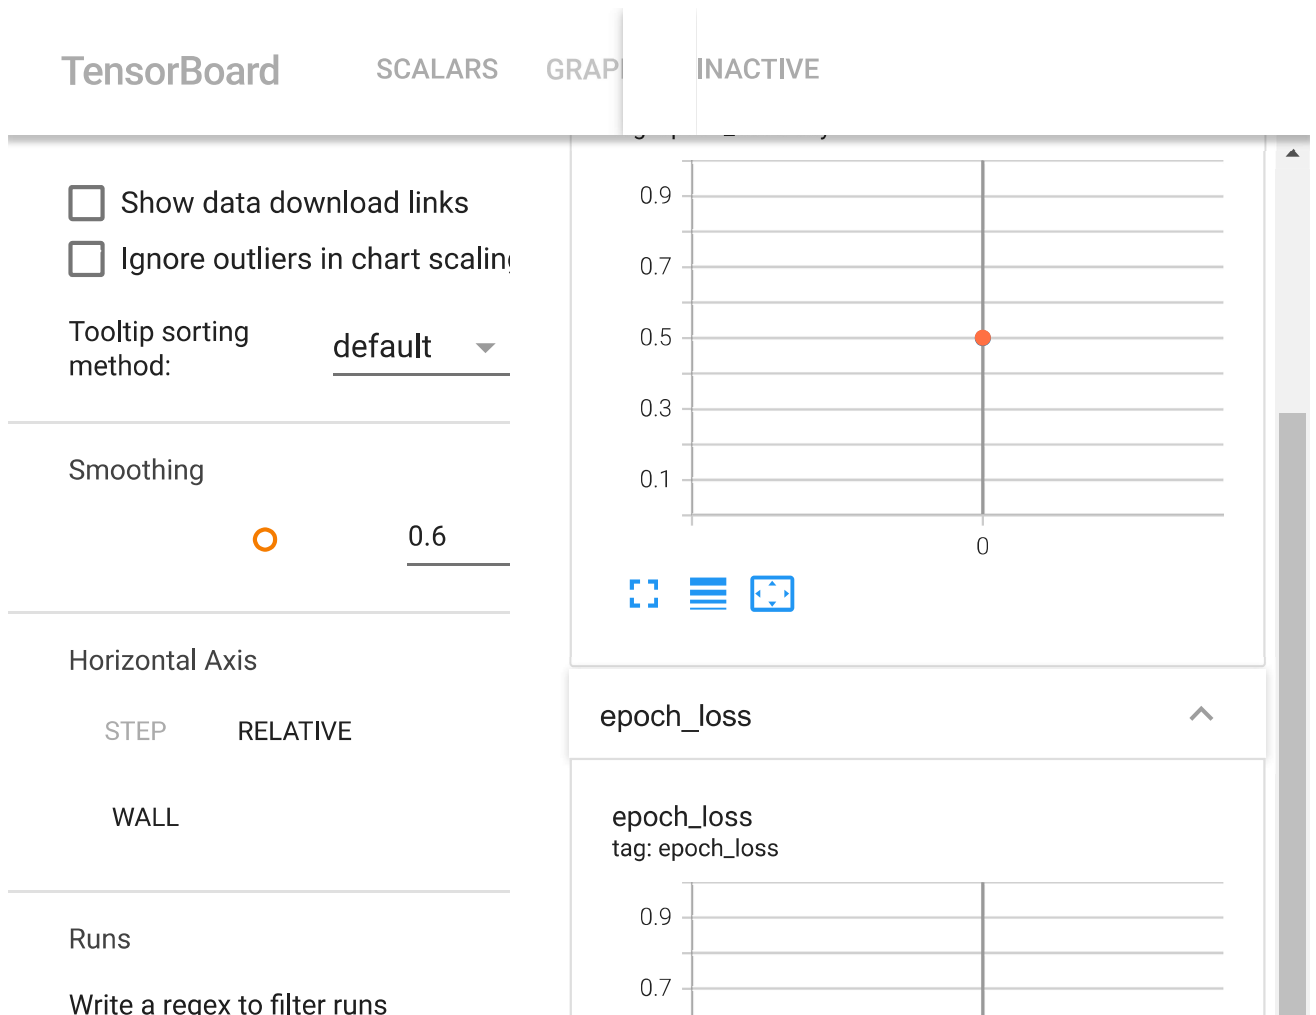
```

Epoch 1/10
 1/1000 [.....] - ETA: 6:02 - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 2/10
1000/1000 [=====] - 1s 1ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 3/10
1000/1000 [=====] - 1s 1ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 4/10
 975/1000 [=====>.] - ETA: 0s - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 1s 1ms/step - loss: 0.0000e+00 - accuracy: 0.0000
<keras.callbacks.History at 0x7fdf6b929210>

```



```
%tensorboard --logdir logs_model3
```



1. Accuracy & AUC are same i.e. 0.5013
2. Training terminated because of invalid loss(NaN) at epoch 0
3. Both Train and test accuracies are same

Model-4

1. Try with any values to get better accuracy/f1 score.

```
#Input layer
input_layer = Input(shape=(2,))
#Dense hidden layer
layer1 = Dense(32,activation='relu',kernel_initializer=tf.keras.initializers.HeNormal(seed=
layer2 = Dense(16,activation='relu',kernel_initializer=tf.keras.initializers.HeNormal(seed=
layer3 = Dense(8,activation='relu',kernel_initializer=tf.keras.initializers.HeNormal(seed=
layer4 = Dense(4,activation='relu',kernel_initializer=tf.keras.initializers.HeNormal(seed=
layer5 = Dense(2,activation='relu',kernel_initializer=tf.keras.initializers.HeNormal(seed=
#output layer
output = Dense(1,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_norm
#Creating a model
```

```

model4 = Model(inputs=input_layer,outputs=output)

optimizer = tf.keras.optimizers.Adam(learning_rate=0.1)

model4.compile(optimizer=optimizer, loss='categorical_crossentropy',metrics=['accuracy'])

#Callbacks
mc=metric_callback(validation_data=[X_test,Y_test])
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
ms = modelSave_callback(filepath)
lr = learning_rate_callback()
tn = terminateNaN_callback()
mcheck = modelcheck_callback()
t_c = tf.keras.callbacks.TensorBoard(log_dir="./logs")

callback_lst = [mc,ms,lr,tn,mcheck,t_c]


model4.fit(X_train,Y_train,epochs=10,validation_data=(X_test,Y_test),batch_size=16,callbacks=callback_lst)

```

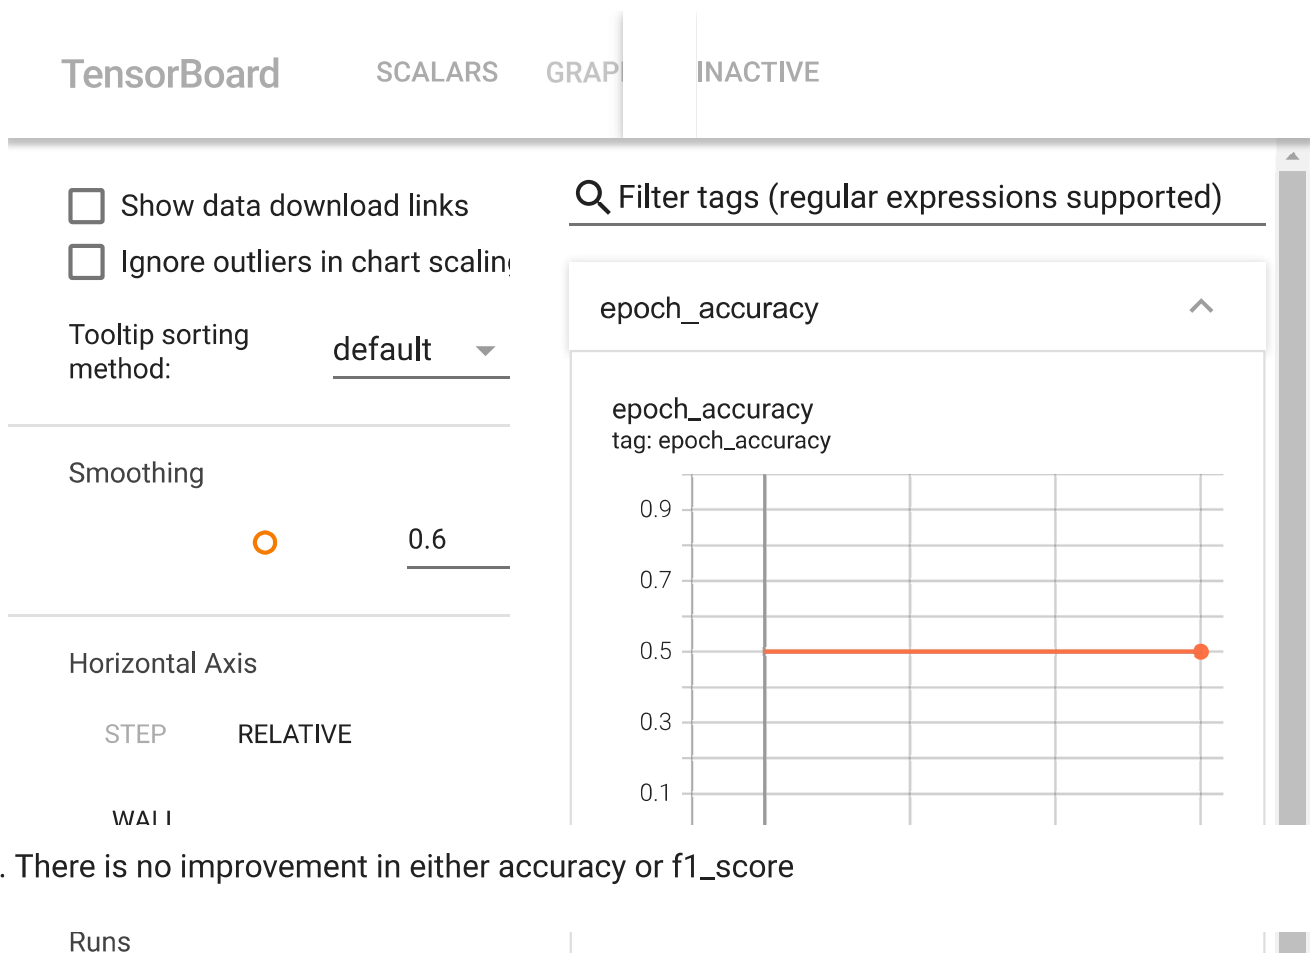
```

Epoch 1/10
 1/1000 [.....] - ETA: 16:23 - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 4s 3ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 2/10
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 3/10
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
Epoch 4/10
 958/1000 [=====>..] - ETA: 0s - loss: 0.0000e+00 - accuracy: 0.0000
1000/1000 [=====] - 2s 2ms/step - loss: 0.0000e+00 - accuracy: 0.0000
<keras.callbacks.History at 0x7fdf6b635490>

```




```
%tensorboard --logdir logs_model4
```



Note

Make sure that you are plotting tensorboard plots either in your notebook or you can try to create a pdf file with all the tensorboard screenshots. Please write your analysis of tensorboard results for each model.



✓ 14s completed at 11:08 PM

×